

project

December 21, 2021

```
[224]: # Importing necessary modules.
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import statsmodels.api as sm
import scipy.stats as stats
from scipy import stats
from sklearn import set_config

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.preprocessing import OrdinalEncoder, LabelEncoder
from sklearn.impute import SimpleImputer

from sklearn.linear_model import LogisticRegression
from sklearn.datasets import make_regression
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline, FeatureUnion
from sklearn.compose import ColumnTransformer
from sklearn.dummy import DummyClassifier
from sklearn.tree import export_graphviz, plot_tree
from IPython.display import Image
from sklearn.metrics import precision_score
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import cross_val_score

from sklearn import preprocessing
from sklearn.metrics import mean_squared_error, make_scorer
import sklearn.metrics as metrics
from sklearn.metrics import r2_score
```

```
from sklearn.metrics import plot_confusion_matrix, classification_report
from sklearn.metrics import confusion_matrix
set_config(display='diagram')
```

0.1 Business Value

Tanzania is currently a 62million population country. And still doesn't have enough well water resources for some of peoples in the country. Water is the basic needs for human body. Tanzania Government is currently working for solve this problem by improving clean water sources. There are many water wells already established, but some of them are non-functional or needs repair.

0.2 Business Problem

In this model, our aim is the predict **functionality** of water points. This will help Tanzania Government for future work. If a water point needs repair or why is not functional and what features affect functionality. With this model, we can help the Tanzanian authorities how to use water sources in a productive way.

```
[225]: # Import and looking the data.
train_data = pd.read_csv('train_data.csv')
test_data = pd.read_csv('test_data.csv')
train_data.head()
```

```
[225]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	\
0	69572	6000.0	2011-03-14	Roman	1390	Roman	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	
4	19728	0.0	2011-07-13	Action In A	0	Artisan	

	longitude	latitude	wpt_name	num_private	...	payment_type	\
0	34.938093	-9.856322	none	0	...	annually	
1	34.698766	-2.147466	Zahanati	0	...	never pay	
2	37.460664	-3.821329	Kwa Mahundi	0	...	per bucket	
3	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	...	never pay	
4	31.130847	-1.825359	Shuleni	0	...	never pay	

	water_quality	quality_group	quantity	quantity_group	\
0	soft	good	enough	enough	
1	soft	good	insufficient	insufficient	
2	soft	good	enough	enough	
3	soft	good	dry	dry	
4	soft	good	seasonal	seasonal	

	source	source_type	source_class	\
0	spring	spring	groundwater	
1	rainwater harvesting	rainwater harvesting	surface	
2	dam	dam	surface	

```

3          machine dbh          borehole  groundwater
4  rainwater harvesting  rainwater harvesting          surface

```

```

          waterpoint_type waterpoint_type_group
0          communal standpipe  communal standpipe
1          communal standpipe  communal standpipe
2  communal standpipe multiple  communal standpipe
3  communal standpipe multiple  communal standpipe
4          communal standpipe  communal standpipe

```

[5 rows x 40 columns]

```

[226]: # Importing labels as our target variable.
labels = pd.read_csv('train_labels.csv')
labels

```

```

[226]:      id  status_group
0    69572    functional
1     8776    functional
2    34310    functional
3    67743  non functional
4    19728    functional
...
59395  60739    functional
59396  27263    functional
59397  37057    functional
59398  31282    functional
59399  26348    functional

```

[59400 rows x 2 columns]

```

[227]: #Looking our data.
train_data.describe()

```

```

[227]:      id  amount_tsh  gps_height  longitude  latitude \
count  59400.000000  59400.000000  59400.000000  59400.000000  5.940000e+04
mean    37115.131768    317.650385    668.297239    34.077427  -5.706033e+00
std    21453.128371    2997.574558    693.116350     6.567432   2.946019e+00
min         0.000000         0.000000   -90.000000     0.000000  -1.164944e+01
25%    18519.750000         0.000000     0.000000    33.090347  -8.540621e+00
50%    37061.500000         0.000000    369.000000    34.908743  -5.021597e+00
75%    55656.500000    20.000000    1319.250000    37.178387  -3.326156e+00
max    74247.000000  350000.000000    2770.000000    40.345193  -2.000000e-08

      num_private  region_code  district_code  population \
count  59400.000000  59400.000000  59400.000000  59400.000000
mean     0.474141    15.297003     5.629747    179.909983

```

std	12.236230	17.587406	9.633649	471.482176
min	0.000000	1.000000	0.000000	0.000000
25%	0.000000	5.000000	2.000000	0.000000
50%	0.000000	12.000000	3.000000	25.000000
75%	0.000000	17.000000	5.000000	215.000000
max	1776.000000	99.000000	80.000000	30500.000000

```

construction_year
count      59400.000000
mean       1300.652475
std        951.620547
min         0.000000
25%         0.000000
50%        1986.000000
75%        2004.000000
max        2013.000000

```

```

[228]: #Looking inside features.(what kind of column with the column names and missing
        ↪ values)
        train_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 40 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    59400 non-null  int64
1   amount_tsh           59400 non-null  float64
2   date_recorded        59400 non-null  object
3   funder                55765 non-null  object
4   gps_height           59400 non-null  int64
5   installer            55745 non-null  object
6   longitude            59400 non-null  float64
7   latitude             59400 non-null  float64
8   wpt_name             59400 non-null  object
9   num_private          59400 non-null  int64
10  basin                59400 non-null  object
11  subvillage           59029 non-null  object
12  region               59400 non-null  object
13  region_code          59400 non-null  int64
14  district_code        59400 non-null  int64
15  lga                  59400 non-null  object
16  ward                 59400 non-null  object
17  population            59400 non-null  int64
18  public_meeting       56066 non-null  object
19  recorded_by          59400 non-null  object
20  scheme_management     55523 non-null  object
21  scheme_name          31234 non-null  object

```

```

22 permit                56344 non-null object
23 construction_year      59400 non-null int64
24 extraction_type         59400 non-null object
25 extraction_type_group   59400 non-null object
26 extraction_type_class   59400 non-null object
27 management              59400 non-null object
28 management_group        59400 non-null object
29 payment                 59400 non-null object
30 payment_type            59400 non-null object
31 water_quality           59400 non-null object
32 quality_group           59400 non-null object
33 quantity                59400 non-null object
34 quantity_group          59400 non-null object
35 source                  59400 non-null object
36 source_type             59400 non-null object
37 source_class            59400 non-null object
38 waterpoint_type         59400 non-null object
39 waterpoint_type_group   59400 non-null object
dtypes: float64(3), int64(7), object(30)
memory usage: 18.1+ MB

```

0.3 Data Understanding

Our data has 39 different columns as feature include 'id'. Some of them has missing values and

We have data from 1960 to 2013 with different funders at 21 different regions in Tanzania. And

We have quantity of the water source with source type as spring, shallow well etc.

We are going to more focus data understanding as looking inside every column.

```

[229]: # Merging target and features.
train_data = train_data.merge(labels,on='id')
train_data

```

```

[229]:      id  amount_tsh  date_recorded      funder  gps_height  \
0    69572      6000.0   2011-03-14      Roman      1390
1     8776         0.0   2013-03-06    Grumeti      1399
2    34310       25.0   2013-02-25  Lottery Club      686
3    67743         0.0   2013-01-28      Unicef      263
4    19728         0.0   2011-07-13  Action In A         0
...    ...      ...      ...      ...      ...
59395  60739       10.0   2013-05-03  Germany Republi      1210
59396  27263      4700.0   2011-05-07    Cefa-njombe      1212
59397  37057         0.0   2011-04-11         NaN         0
59398  31282         0.0   2011-03-08      Malec         0
59399  26348         0.0   2011-03-23    World Bank      191

```

	installer	longitude	latitude	wpt_name	num_private	\
0	Roman	34.938093	-9.856322	none	0	
1	GRUMETI	34.698766	-2.147466	Zahanati	0	
2	World vision	37.460664	-3.821329	Kwa Mahundi	0	
3	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	
4	Artisan	31.130847	-1.825359	Shuleni	0	
...	
59395	CES	37.169807	-3.253847	Area Three Namba 27	0	
59396	Cefa	35.249991	-9.070629	Kwa Yahona Kuvala	0	
59397	NaN	34.017087	-8.750434	Mashine	0	
59398	Musa	35.861315	-6.378573	Mshoro	0	
59399	World	38.104048	-6.747464	Kwa Mzee Lugawa	0	

	...	water_quality	quality_group	quantity	quantity_group	\
0	...	soft	good	enough	enough	
1	...	soft	good	insufficient	insufficient	
2	...	soft	good	enough	enough	
3	...	soft	good	dry	dry	
4	...	soft	good	seasonal	seasonal	
...	
59395	...	soft	good	enough	enough	
59396	...	soft	good	enough	enough	
59397	...	fluoride	fluoride	enough	enough	
59398	...	soft	good	insufficient	insufficient	
59399	...	salty	salty	enough	enough	

		source	source_type	source_class	\
0		spring	spring	groundwater	
1	rainwater	harvesting	rainwater harvesting	surface	
2		dam	dam	surface	
3		machine dbh	borehole	groundwater	
4	rainwater	harvesting	rainwater harvesting	surface	
...		
59395		spring	spring	groundwater	
59396		river	river/lake	surface	
59397		machine dbh	borehole	groundwater	
59398		shallow well	shallow well	groundwater	
59399		shallow well	shallow well	groundwater	

		waterpoint_type	waterpoint_type_group	status_group
0		communal standpipe	communal standpipe	functional
1		communal standpipe	communal standpipe	functional
2	communal	standpipe multiple	communal standpipe	functional
3	communal	standpipe multiple	communal standpipe	non functional
4		communal standpipe	communal standpipe	functional
...	
59395		communal standpipe	communal standpipe	functional

59396	communal	standpipe	communal	standpipe	functional
59397		hand pump		hand pump	functional
59398		hand pump		hand pump	functional
59399		hand pump		hand pump	functional

[59400 rows x 41 columns]

0.4 Cleaning Process

0.4.1 Train Data Cleaning

```
[230]: #Looking for missing values.
train_data.isna().sum()
```

```
[230]: id                0
amount_tsh              0
date_recorded           0
funder                 3635
gps_height              0
installer              3655
longitude               0
latitude                0
wpt_name                0
num_private             0
basin                  0
subvillage              371
region                  0
region_code             0
district_code           0
lga                     0
ward                    0
population              0
public_meeting          3334
recorded_by             0
scheme_management       3877
scheme_name             28166
permit                  3056
construction_year       0
extraction_type         0
extraction_type_group   0
extraction_type_class   0
management              0
management_group        0
payment                 0
payment_type            0
water_quality           0
quality_group           0
```

```

quantity                0
quantity_group          0
source                  0
source_type             0
source_class            0
waterpoint_type        0
waterpoint_type_group  0
status_group           0
dtype: int64

```

```

[231]: # Turning target column from ternary to binary.
train_data['status_group'] = train_data['status_group'].replace('functional_
↳needs repair','non functional')

```

```

[232]: # Checking target variable values.
train_data['status_group'].value_counts()

```

```

[232]: functional          32259
non functional          27141
Name: status_group, dtype: int64

```

```

[233]: # Filling some columns as feature engineering.
train_data['funder'].fillna('Unkown',inplace=True)
train_data['funder'].replace(to_replace = '0', value ='Unknown' , inplace=True)

train_data['installer'].fillna('Unkown',inplace=True)

train_data['subvillage'].fillna('Missing',inplace=True)
train_data['public_meeting'].fillna(False,inplace=True)

```

```

[234]: # Looking inside scheme management column.
train_data['scheme_management'].value_counts()

```

```

[234]: VWC                36793
WUG                5206
Water authority    3153
WUA                2883
Water Board       2748
Parastatal        1680
Private operator   1063
Company           1061
Other              766
SWC                97
Trust              72
None               1
Name: scheme_management, dtype: int64

```



```
[235]: # Filling scheme management column.
train_data['scheme_management'].fillna('Missing', inplace=True)
```

```
[236]: # Filling permit column.
train_data['permit'].fillna(False, inplace=True)
```

```
[237]: # Boolean column converting process.
le = preprocessing.LabelEncoder()
train_data['public_meeting'] = le.fit_transform(train_data.public_meeting.
→values)
train_data['permit'] = le.fit_transform(train_data.permit.values)
train_data.dtypes
```

```
[237]: id                int64
amount_tsh              float64
date_recorded           object
funder                  object
gps_height              int64
installer               object
longitude               float64
latitude                float64
wpt_name                object
num_private              int64
basin                   object
subvillage              object
region                  object
region_code             int64
district_code           int64
lga                     object
ward                    object
population              int64
public_meeting          int64
recorded_by             object
scheme_management       object
scheme_name             object
permit                  int64
construction_year       int64
extraction_type          object
extraction_type_group   object
extraction_type_class   object
management              object
management_group        object
payment                 object
payment_type            object
water_quality           object
quality_group           object
quantity                object
```

quantity_group	object
source	object
source_type	object
source_class	object
waterpoint_type	object
waterpoint_type_group	object
status_group	object
dtype:	object

0.4.2 Looking Inside Each Column

Column Names

amount_tsh - Total static head (amount water available to waterpoint)
 date_recorded - The date the row was entered
 funder - Who funded the well
 gps_height - Altitude of the well
 installer - Organization that installed the well
 longitude - GPS coordinate
 latitude - GPS coordinate
 wpt_name - Name of the waterpoint if there is one
 num_private -
 basin - Geographic water basin
 subvillage - Geographic location
 region - Geographic location
 region_code - Geographic location (coded)
 district_code - Geographic location (coded)
 lga - Geographic location
 ward - Geographic location
 population - Population around the well
 public_meeting - True/False
 recorded_by - Group entering this row of data
 scheme_management - Who operates the waterpoint
 scheme_name - Who operates the waterpoint
 permit - If the waterpoint is permitted
 construction_year - Year the waterpoint was constructed
 extraction_type - The kind of extraction the waterpoint uses
 extraction_type_group - The kind of extraction the waterpoint uses
 extraction_type_class - The kind of extraction the waterpoint uses
 management - How the waterpoint is managed
 management_group - How the waterpoint is managed
 payment - What the water costs
 payment_type - What the water costs
 water_quality - The quality of the water
 quality_group - The quality of the water
 quantity - The quantity of water
 quantity_group - The quantity of water
 source - The source of the water

source_type - The source of the water
source_class - The source of the water
waterpoint_type - The kind of waterpoint
waterpoint_type_group - The kind of waterpoint

Amount_tsh Column

```
[238]: train_data['amount_tsh'].value_counts()
```

```
[238]: 0.0          41639
      500.0       3102
      50.0       2472
      1000.0     1488
      20.0       1463
      ...
      8500.0        1
      6300.0        1
      220.0         1
      138000.0      1
      12.0          1
      Name: amount_tsh, Length: 98, dtype: int64
```


```
[239]: # Defining function for feature engineering on amount column.
      def split_amount(amount):
          if amount < 10:
              return 'Less amount water source at this point.'
          return 'Enough amount water source at this point.'
```

```
[240]: #Applying function to this column.
      train_data['amount_tsh'] = train_data['amount_tsh'].apply(split_amount)
```

```
[241]: train_data['amount_tsh'].value_counts()
```

```
[241]: Less amount water source at this point.    42295
      Enough amount water source at this point.  17105
      Name: amount_tsh, dtype: int64
```

```
[242]: # Creating visualizing for this column with functionality.
      sns.set_theme(style="darkgrid")
      fig, ax = plt.subplots(figsize=(10,8))
      ax = sns.countplot(x='amount_tsh', hue="status_group", data=train_data )
```



If water point has enough source , there seems to be likely functional , on the other hand if there is not enough water source at the point almost half of them functional.

Date Recorded Column

```
[243]: # Feature engineering in date recorded column.
```

```
years = []
for i in train_data.date_recorded:
    years.append(i[:4])
```

```
[244]: train_data['date_recorded'] = years
```

```
[245]: # For date column we are going to use just 'years'.
```

```
[246]: train_data['date_recorded'].value_counts()
```

```
[246]: 2011    28674
      2013    24271
      2012     6424
      2004      30
      2002       1
```

Name: date_recorded, dtype: int64

```
[247]: # Replacing outliers with other less value.  
train_data['date_recorded'].replace('2002', '2012', inplace=True)  
train_data['date_recorded'].replace('2004', '2012', inplace=True)
```

```
[248]: train_data['date_recorded'].value_counts()
```

```
[248]: 2011      28674  
      2013      24271  
      2012       6455  
      Name: date_recorded, dtype: int64
```

Funder Column

```
[249]: train_data.funder.value_counts()
```

```
[249]: Government Of Tanzania      9084  
      Unkown                    3635  
      Danida                   3114  
      Hesawa                   2202  
      Rwssp                    1374  
      ...  
      Tanedaps Society          1  
      Dwe/ubalozi Wa Marekani    1  
      Noeli Mahobokela           1  
      Deogratus Kasima           1  
      Muslimehefen International 1  
      Name: funder, Length: 1897, dtype: int64
```

```
[250]: funders = train_data['funder']
```

```
[251]: # Changing values at the funder column.  
train_data.loc[train_data['funder']  
               .value_counts()  
               [train_data['funder']  
               .values < 201, 'funder']] = "Others"
```

```
[252]: train_data['funder'].value_counts()
```

```
[252]: Others      18842  
      Government Of Tanzania      9084  
      Unkown                    3635  
      Danida                   3114  
      Hesawa                   2202  
      Rwssp                    1374  
      World Bank                1349  
      Kkkt                     1287
```

World Vision	1246
Unicef	1057
Tasaf	877
District Council	843
Dhv	829
Private Individual	826
Dwsp	811
Unknown	781
Norad	765
Germany Republi	610
Tcrs	602
Ministry Of Water	590
Water	583
Dwe	484
Netherlands	470
Hifab	450
Adb	448
Lga	442
Amref	425
Fini Water	393
Oxfam	359
Wateraid	333
Rc Church	321
Isf	316
Rudep	312
Mission	301
Private	295
Jaica	280
Roman	275
Rural Water Supply And Sanitat	270
Adra	263
Ces(gmbh)	260
Jica	259
Shipo	241
Wsdp	234
Rc	230
Finw	219
Dh	213

Name: funder, dtype: int64

Gps_height Column


```
[253]: # Creating function for gps column seperation.
def split_gps(gps):
    if gps < 200:
        return 'Low altitude water source'
    if gps >= 200:
        return 'High altitude water source'
```

```
[254]: train_data['gps_height'] = train_data['gps_height'].apply(split_gps)
```

```
[255]: # Checking this column values.  
train_data.gps_height.value_counts(normalize=True)
```

```
[255]: High altitude water source    0.569579  
Low altitude water source        0.430421  
Name: gps_height, dtype: float64
```

```
[256]: # Craeting graph for gps height with functionality.  
fig, ax = plt.subplots(figsize=(10,8))  
ax = sns.countplot(x='gps_height', hue="status_group", data=train_data )
```



It tend to be at the high altitude more functional likely. At the low altitude looks like even functional and non functional water points.

Installer Column

```
[257]: train_data.installer.value_counts()
```

```
[257]: DWE                17402
      Unkown            3655
      Government       1825
      RWE               1206
      Commu             1060
      ...
      MSIKITI           1
      LEI               1
      Busoga trust      1
      Lualu Kaima       1
      Government /SDA   1
      Name: installer, Length: 2146, dtype: int64
```

```
[258]: # Changing values at the intaller column.
      train_data.loc[train_data['installer']
                     .value_counts()
                     [train_data['installer']]
                     .values > 1001, 'installer'] = "Above 1000 water point_
      ↪installers."
```

```
[259]: # Changing values at the installer column.
      train_data.loc[train_data['installer']
                     .value_counts()
                     [train_data['installer']]
                     .values < 100, 'installer'] = "Under 100 water point installers."
```

```
[260]: # Changing values at the installer column.
      train_data.loc[train_data['installer']
                     .value_counts()
                     [train_data['installer']]
                     .values < 1000, 'installer'] = "Under 1000 water point_
      ↪installers."
```

```
[261]: # Checking installer column values.
      train_data.installer.value_counts()
```

```
[261]: Above 1000 water point installers.    26198
      Under 1000 water point installers.    19657
      Under 100 water point installers.     13545
      Name: installer, dtype: int64
```


Basin Column

```
[262]: train_data.basin.value_counts()
```



```
[262]: Lake Victoria      10248
      Pangani             8940
      Rufiji              7976
      Internal            7785
      Lake Tanganyika     6432
      Wami / Ruvu         5987
      Lake Nyasa          5085
      Ruvuma / Southern Coast 4493
      Lake Rukwa          2454
      Name: basin, dtype: int64
```

```
[263]: # Creating graph for basin column with functionality.
fig, ax = plt.subplots(figsize=(16,8))
ax = sns.countplot(x='basin', hue="status_group", data=train_data)
```



Subvillage Column

```
[264]: train_data['subvillage'].value_counts()
```

```
[264]: Madukani      508
      Shuleni       506
      Majengo       502
      Kati          373
      Missing       371
      ...
      Matakuja      1
      Nyamuko        1
      Salaliya       1
      Hongi Juu      1
```

Wami 1
 Name: subvillage, Length: 19288, dtype: int64

[265]: *# Feature engineering on subvillage column.*

```
train_data.loc[train_data['subvillage']
               .value_counts()
               [train_data['subvillage']]
               .values > 100]
```

[265]:

	id	amount_tsh	date_recorded	\
2	34310	Enough amount water source at this point.	2013	
20	48375	Enough amount water source at this point.	2011	
26	55012	Enough amount water source at this point.	2013	
42	52019	Enough amount water source at this point.	2011	
70	21990	Enough amount water source at this point.	2011	
...	
59358	44951	Enough amount water source at this point.	2011	
59365	8810	Less amount water source at this point.	2011	
59369	47527	Less amount water source at this point.	2011	
59384	72148	Less amount water source at this point.	2011	
59385	34473	Enough amount water source at this point.	2012	

	funder	gps_height	\
2	Others	High altitude water source	
20	Others	High altitude water source	
26	Others	High altitude water source	
42	Others	High altitude water source	
70	Government Of Tanzania	High altitude water source	
...	
59358	Unicef	High altitude water source	
59365	Unicef	High altitude water source	
59369	Rc Church	Low altitude water source	
59384	Others	Low altitude water source	
59385	Jaica	High altitude water source	

	installer	longitude	latitude	\
2	Under 1000 water point installers.	37.460664	-3.821329	
20	Under 1000 water point installers.	34.473430	-9.594990	
26	Under 100 water point installers.	39.370777	-9.942532	
42	Under 100 water point installers.	34.814574	-9.032503	
70	Above 1000 water point installers.	35.818981	-8.934950	
...	
59358	Above 1000 water point installers.	34.631938	-8.723208	
59365	Above 1000 water point installers.	34.594790	-9.072904	
59369	Under 100 water point installers.	33.670049	-9.001535	
59384	Under 100 water point installers.	30.667805	-2.483710	
59385	Under 100 water point installers.	33.951681	-2.021854	

	wpt_name	num_private	...	water_quality	quality_group	\
2	Kwa Mahundi	0	...	soft	good	
20	none	0	...	soft	good	
26	Ruhoma Primary School	0	...	soft	good	
42	Zahanati-Misssion	0	...	soft	good	
70	Kwampalanji	0	...	unknown	unknown	
...	
59358	Kwa Helena Mabena	0	...	soft	good	
59365	Kwa Yohane Mhanza	0	...	soft	good	
59369	Kwa Paval Dinno	0	...	soft	good	
59384	Chamkube	0	...	soft	good	
59385	Kwa Marunda	0	...	salty	salty	

	quantity	quantity_group	source	source_type	source_class	\
2	enough	enough	dam	dam	surface	
20	enough	enough	spring	spring	groundwater	
26	enough	enough	machine dbh	borehole	groundwater	
42	enough	enough	spring	spring	groundwater	
70	insufficient	insufficient	shallow well	shallow well	groundwater	
...	
59358	enough	enough	river	river/lake	surface	
59365	enough	enough	river	river/lake	surface	
59369	enough	enough	river	river/lake	surface	
59384	insufficient	insufficient	spring	spring	groundwater	
59385	enough	enough	machine dbh	borehole	groundwater	

	waterpoint_type	waterpoint_type_group	status_group
2	communal standpipe multiple	communal standpipe	functional
20	communal standpipe	communal standpipe	functional
26	hand pump	hand pump	functional
42	communal standpipe	communal standpipe	functional
70	hand pump	hand pump	non functional
...
59358	communal standpipe	communal standpipe	functional
59365	communal standpipe	communal standpipe	non functional
59369	communal standpipe	communal standpipe	functional
59384	communal standpipe	communal standpipe	non functional
59385	hand pump	hand pump	functional

[4841 rows x 41 columns]

Region Column

```
[283]: train_data['region'].value_counts()
```

```
[283]: Iringa          5294
       Shinyanga       4982
```

Mbeya	4639
Kilimanjaro	4379
Morogoro	4006
Arusha	3350
Kagera	3316
Mwanza	3102
Kigoma	2816
Ruvuma	2640
Pwani	2635
Tanga	2547
Dodoma	2201
Singida	2093
Mara	1969
Tabora	1959
Rukwa	1808
Mtwara	1730
Manyara	1583
Lindi	1546
Dar es Salaam	805

Name: region, dtype: int64

```
[310]: sorted_region = train_data['population'].groupby(train_data['region']).sum().
        ↪sort_values()
```

```
[317]: train_data['population'].groupby(train_data['region']).sum()
```

```
[317]: region
Arusha      878782
Dar es Salaam  193879
Dodoma      618481
Iringa      826331
Kagera      931796
Kigoma     1417392
Kilimanjaro  463070
Lindi       563370
Manyara     503043
Mara       1060886
Mbeya      1303559
Morogoro    1060090
Mtwara      462674
Mwanza      971145
Pwani       921177
Rukwa       674566
Ruvuma      656638
Shinyanga   1424109
Singida     584765
Tabora      550479
```

Tanga 628482
 Name: population, dtype: int64

[321]: grouped_region

[321]:

	region	id	longitude	latitude \
17	Shinyanga	182679138	132,195.96	-13,902.38
5	Kigoma	102667532	85,084.53	-12,098.15
10	Mbeya	174055235	155,552.65	-42,202.92
9	Mara	74861850	67,249.11	-3,425.87
11	Morogoro	148280588	148,405.91	-29,690.72
13	Mwanza	116488810	75,945.14	-6,014.05
4	Kagera	123713407	103,569.50	-6,504.22
14	Pwani	97304169	102,445.97	-18,471.23
0	Arusha	125452610	122,451.59	-10,875.62
3	Iringa	194159068	184,739.36	-47,157.36
15	Rukwa	67564706	56,562.17	-13,305.58
16	Ruvuma	99180138	94,325.36	-28,444.36
20	Tanga	94829204	98,076.07	-12,915.49
2	Dodoma	80760619	79,333.22	-13,049.14
18	Singida	76474025	72,714.67	-10,265.46
7	Lindi	57574121	60,289.53	-15,113.21
19	Tabora	72555411	64,415.02	-9,251.57
8	Manyara	59102904	56,881.80	-6,785.07
6	Kilimanjaro	163126393	164,234.87	-15,426.44
12	Mtwara	63791546	68,157.98	-18,477.19
1	Dar es Salaam	30017353	31,568.72	-5,562.29

	num_private	region_code	district_code	population	public_meeting \
17	0	84598	17785	1424109	3761
5	0	45056	5597	1417392	2663
10	0	55668	19941	1303559	3914
9	73	39380	6031	1060886	794
11	426	20030	12141	1060090	3861
13	0	58828	12550	971145	2695
4	0	59688	28869	931796	3277
14	3190	71162	52289	921177	2234
0	32	13872	22636	878782	3075
3	63	58234	18893	826331	5012
15	0	27120	3725	674566	1419
16	905	26400	7931	656638	1897
20	16298	10222	9211	628482	2362
2	0	2201	6493	618481	2192
18	137	27209	4373	584765	1908
7	10	101584	41540	563370	1072
19	0	27426	5583	550479	1397
8	57	33243	4023	503043	1506

6	6520	13137	15956	463070	3961
12	453	127917	37108	462674	1672
1	0	5635	1732	193879	339

	permit	construction_year
17	2364	328430
5	1656	5606470
10	2134	0
9	998	3899032
11	3967	7960149
13	2975	707888
4	2521	0
14	1530	4925027
0	2650	6594069
3	3174	9831406
15	1162	3593685
16	1798	5241299
20	1378	4887905
2	989	0
18	805	4158195
7	813	2858609
19	1325	0
8	1506	3143319
6	3796	8672984
12	1311	3273470
1	0	1576820

```
[323]: grouped_region = train_data.groupby('region').sum().reset_index()
grouped_region.sort_values(by='population',ascending=True,inplace=True)
```

```
[365]: #Creating function to fix scientific notations.
def notation(x, pos):
    """The two args are the value and tick position"""
    if x >= 1e6:
        s = '{:1.1f}M'.format(x*1e-6)
    else:
        s = '{:1.0f}K'.format(x*1e-2)
    return s
```


```
[366]: #Creating visual for region populations.
colors=[]
for region in grouped_region['region']:
    if (
        (region=='Shinyanga') or
        (region=='Kigoma') or
        (region=='Mbeya') or
        (region=='Mara')
```

```

):
    colors.append('royalblue')
else:
    colors.append('slategray')

fig, ax = plt.subplots(figsize=(12,8))
ax.barh(grouped_region['region'],grouped_region['population'],color=colors)
ax.set_title('Most populated regions',fontsize=20)
plt.xlabel('Population',fontsize=20)
plt.ylabel('Regions',fontsize=20)
ax.xaxis.set_major_formatter(notation)
plt.style.use('ggplot')


```



```

[347]: # Creating graph for regions with functionality.
fig, ax = plt.subplots(figsize=(8,12))
ax = sns.countplot(y='region',
                  hue="status_group",
                  data=train_data,order=train_data.region.value_counts().index
                  )

```



Seems to be most functional water points at 'Iringa', 'Shinyanga' and 'Kilimanjaro' region in order. What about these regions populations ?

Region_Code

```
[274]: train_data.region_code.value_counts()
```



```
[274]: 11    5300
      17    5011
      12    4639
      3    4379
      5    4040
      18    3324
      19    3047
      2    3024
      16    2816
      10    2640
      4    2513
      1    2201
      13    2093
      14    1979
      20    1969
      15    1808
      6    1609
      21    1583
      80    1238
      60    1025
      90     917
      7     805
      99     423
      9     390
      24     326
      8     301
      Name: region_code, dtype: int64
```

```
[275]: train_data['region_code'].replace(40,8,inplace=True)
```

```
[276]: train_data.region_code.value_counts()
```

```
[276]: 11    5300
      17    5011
      12    4639
      3    4379
      5    4040
      18    3324
      19    3047
      2    3024
      16    2816
      10    2640
      4    2513
      1    2201
      13    2093
      14    1979
      20    1969
```

```

15    1808
6     1609
21    1583
80    1238
60    1025
90     917
7      805
99     423
9      390
24     326
8      301
Name: region_code, dtype: int64

```

```
[277]: train_data.district_code.value_counts()
```

```

[277]: 1     12203
2     11173
3      9998
4      8999
5      4356
6      4074
7      3343
8      1043
30       995
33       874
53       745
43       505
13       391
23       293
63       195
62       109
60        63
0         23
80         12
67          6
Name: district_code, dtype: int64

```

Region code and district code columns almost represent same things. Also they are in multicollinearity with each other. We are going to drop district code column.

Population

```
[278]: train_data.population.value_counts()
```

```

[278]: 0     21381
1      7025
200     1940
150     1892

```

```

250      1681
      ...
3241      1
1960      1
1685      1
2248      1
1439      1
Name: population, Length: 1049, dtype: int64

```

```

[279]: #train_data['population'].replace(to_replace = 0 , value =281, inplace=True)
      ↪#changing 0's to mean

```

```

[388]: def split_population(p):
      if p < 10:
          return 'Low population at the water source.'
      if 10 < p < 200:
          return 'Under 200 people population.'
      return 'High population.'

```

```

[389]: train_data['population'] = train_data['population'].apply(split_population)

```

```

[390]: train_data.population.value_counts()

```

```

[390]: High population.                38709
      Under 200 people population.    13545
      Low population at the water source.  7146
      Name: population, dtype: int64

```

Public_meeting

```

[53]: train_data.public_meeting.value_counts(normalize=True)

```

```

[53]: 1    0.858771
      0    0.141229
      Name: public_meeting, dtype: float64

```

Scheme_management Column

```

[54]: train_data.scheme_management.value_counts()

```

```

[54]: VWC                36793
      WUG                5206
      Missing           3877
      Water authority    3153
      WUA               2883
      Water Board       2748
      Parastatal        1680

```

Private operator	1063
Company	1061
Other	766
SWC	97
Trust	72
None	1

Name: scheme_management, dtype: int64

```
[55]: train_data.loc[train_data['scheme_management']
        .value_counts()
        [train_data['scheme_management']]
        .values < 200, 'scheme_management'] = "Other"
```

```
[56]: train_data.scheme_management.value_counts()
```

VWC	36793
WUG	5206
Missing	3877
Water authority	3153
WUA	2883
Water Board	2748
Parastatal	1680
Private operator	1063
Company	1061
Other	936

Name: scheme_management, dtype: int64

Construction_year Column

```
[57]: train_data.construction_year.value_counts(normalize =True)
```

0	0.348636
2010	0.044529
2008	0.043990
2009	0.042643
2000	0.035202
2007	0.026717
2006	0.024764
2003	0.021650
2011	0.021145
2004	0.018906
2012	0.018249
2002	0.018098
1978	0.017458
1995	0.017071
2005	0.017020
1999	0.016481
1998	0.016263

```

1990    0.016061
1985    0.015909
1980    0.013653
1996    0.013653
1984    0.013114
1982    0.012525
1994    0.012424
1972    0.011919
1974    0.011380
1997    0.010842
1992    0.010774
1993    0.010236
2001    0.009091
1988    0.008771
1983    0.008215
1975    0.007357
1986    0.007306
1976    0.006970
1970    0.006919
1991    0.005455
1989    0.005320
1987    0.005084
1981    0.004007
1977    0.003401
1979    0.003232
1973    0.003098
2013    0.002963
1971    0.002441
1960    0.001717
1967    0.001481
1963    0.001431
1968    0.001296
1969    0.000993
1964    0.000673
1962    0.000505
1961    0.000354
1965    0.000320
1966    0.000286
Name: construction_year, dtype: float64

```

```
[58]: train_data[train_data.construction_year>0]['construction_year'].median()
```

```
[58]: 2000.0
```

```
[59]: # Since construction_year column's median 2000 , We are going to replace
      ↪missing value 0's to 2000.
```

```
train_data['construction_year'].replace(to_replace = 0, value = 2000,  
↳ inplace=True)
```

```
[60]: train_data.construction_year.value_counts()
```

```
[60]: 2000      22800  
      2010      2645  
      2008      2613  
      2009      2533  
      2007      1587  
      2006      1471  
      2003      1286  
      2011      1256  
      2004      1123  
      2012      1084  
      2002      1075  
      1978      1037  
      1995      1014  
      2005      1011  
      1999       979  
      1998       966  
      1990       954  
      1985       945  
      1996       811  
      1980       811  
      1984       779  
      1982       744  
      1994       738  
      1972       708  
      1974       676  
      1997       644  
      1992       640  
      1993       608  
      2001       540  
      1988       521  
      1983       488  
      1975       437  
      1986       434  
      1976       414  
      1970       411  
      1991       324  
      1989       316  
      1987       302  
      1981       238  
      1977       202  
      1979       192  
      1973       184
```

2013	176
1971	145
1960	102
1967	88
1963	85
1968	77
1969	59
1964	40
1962	30
1961	21
1965	19
1966	17

Name: construction_year, dtype: int64


```
[61]: def split_year(y):
        if y < 1970:
            return '60`s year construction.'
        if y < 1980:
            return '70`s year construction.'
        if y < 1990:
            return '80`s year construction.'
        if y < 2000:
            return '90`s year construction.'
        if y < 2010:
            return '2000`s year construction.'
        if y < 2020:
            return '2010`s year construction.'
```

```
[62]: train_data['construction_year'] = train_data['construction_year'].
        ↪apply(split_year)
```

```
[63]: train_data['construction_year'].value_counts()
```

```
[63]: 2000`s year construction.    36039
        90`s year construction.    7678
        80`s year construction.    5578
        2010`s year construction.  5161
        70`s year construction.    4406
        60`s year construction.    538
        Name: construction_year, dtype: int64
```

```
[64]: fig, ax = plt.subplots(figsize=(14,8))
        ax = sns.countplot(x='construction_year', hue="status_group", data=train_data )
```



Looks like most of the water points before 90's construction non functional and needs repair.

Extraction_type Column

```
[65]: train_data['extraction_type'].value_counts()
```

```
[65]: gravity                26780
      nira/tanira             8154
      other                   6430
      submersible             4764
      sw n 80                  3670
      mono                    2865
      india mark ii           2400
      afridev                  1770
      ksb                      1415
      other - rope pump        451
      other - sw n 81          229
      windmill                 117
      india mark iii           98
      cemo                      90
      other - play pump        85
      walimi                    48
      climax                   32
      other - mkulima/shinyanga 2
      Name: extraction_type, dtype: int64
```



```
[66]: #Reducing this column's features.
train_data.loc[train_data['extraction_type']
               .value_counts()
               [train_data['extraction_type']]
               .values < 500, 'extraction_type'] = "Other"
```

```
[67]: train_data['extraction_type'].value_counts()
```

```
[67]: gravity          26780
      nira/tanira      8154
      other           6430
      submersible     4764
      swm 80          3670
      mono            2865
      india mark ii   2400
      afridev         1770
      ksb             1415
      Other           1152
      Name: extraction_type, dtype: int64
```

```
[68]: train_data['extraction_type_group'].value_counts()
```

```
[68]: gravity          26780
      nira/tanira      8154
      other           6430
      submersible     6179
      swm 80          3670
      mono            2865
      india mark ii   2400
      afridev         1770
      rope pump        451
      other handpump   364
      other motorpump  122
      wind-powered     117
      india mark iii    98
      Name: extraction_type_group, dtype: int64
```

```
[69]: train_data['extraction_type_class'].value_counts()
```

```
[69]: gravity          26780
      handpump         16456
      other           6430
      submersible     6179
      motorpump        2987
      rope pump        451
      wind-powered     117
      Name: extraction_type_class, dtype: int64
```

```
[70]: # Replacing this column's unneceseray values.
train_data['extraction_type_class'].replace('rope pump','other',inplace=True)
train_data['extraction_type_class'].replace('wind-powered','other',inplace=True)
```

Management Column

```
[71]: train_data['management'].value_counts()
```

```
[71]: vwc                40507
      wug                6515
      water board       2933
      wua                2535
      private operator   1971
      parastatal         1768
      water authority     904
      other              844
      company            685
      unknown            561
      other - school      99
      trust              78
      Name: management, dtype: int64
```


```
[72]: #Reducing this column's features.
train_data.loc[train_data['management']
               .value_counts()
               [train_data['management']]
               .values < 900, 'management'] = "Other"
```

```
[73]: train_data['management'].replace('parastatal','water authority',inplace=True)
```

```
[74]: train_data['management'].value_counts()
```

```
[74]: vwc                40507
      wug                6515
      water board       2933
      water authority    2672
      wua                2535
      Other             2267
      private operator   1971
      Name: management, dtype: int64
```

```
[75]: # Plotting management with functionality.
fig, ax = plt.subplots(figsize=(14,8))
ax = sns.countplot(x='management', hue="status_group", data=train_data )
```



```
[76]: train_data['management_group'].value_counts()
```

```
[76]: user-group      52490
      commercial      3638
      parastatal      1768
      other           943
      unknown         561
      Name: management_group, dtype: int64
```

```
[77]: train_data['management_group'].replace('unkown', 'other', inplace=True)
```

Payment and Other Columns

```
[78]: train_data['payment'].value_counts()
```

```
[78]: never pay          25348
      pay per bucket    8985
      pay monthly       8300
      unknown           8157
      pay when scheme fails 3914
      pay annually       3642
      other             1054
      Name: payment, dtype: int64
```

```
[79]: train_data['payment_type'].value_counts()
```

```
[79]: never pay      25348
      per bucket    8985
      monthly      8300
      unknown       8157
      on failure    3914
      annually     3642
      other         1054
      Name: payment_type, dtype: int64
```

```
[80]: train_data['water_quality'].value_counts()
```

```
[80]: soft          50818
      salty         4856
      unknown       1876
      milky          804
      coloured       490
      salty abandoned 339
      fluoride        200
      fluoride abandoned 17
      Name: water_quality, dtype: int64
```

```
[81]: train_data.loc[train_data['water_quality']
                    .value_counts()
                    [train_data['water_quality']]
                    .values < 9000, 'water_quality'] = "Bad"
```

```
[82]: train_data['water_quality'].value_counts()
```

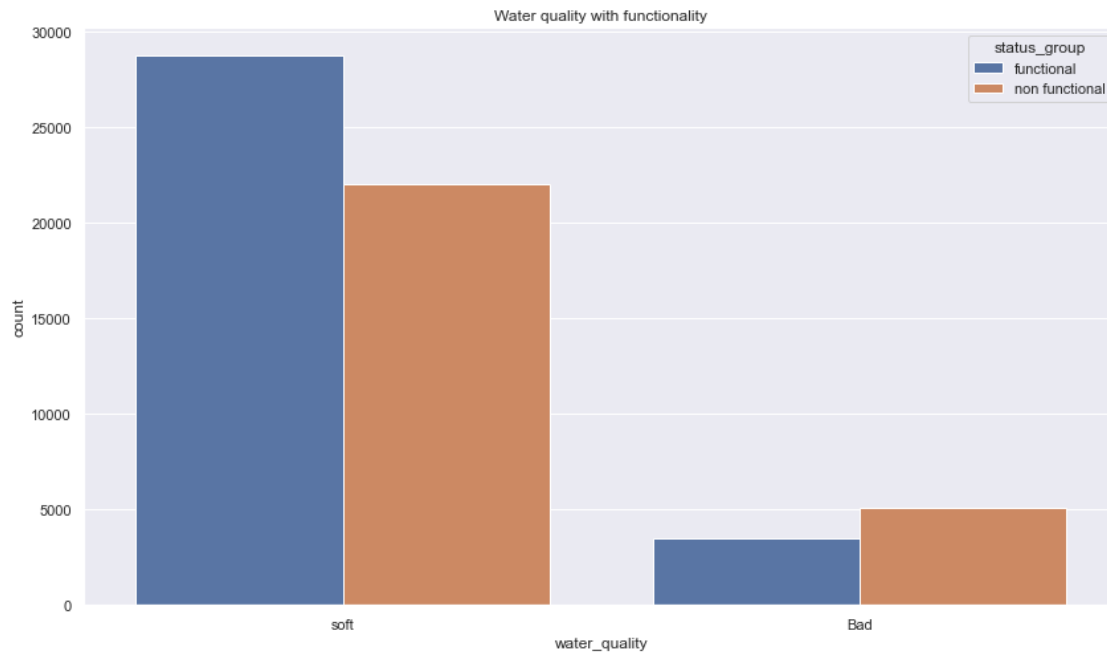
```
[82]: soft      50818
      Bad       8582
      Name: water_quality, dtype: int64
```

```
[83]: train_data['quality_group'].value_counts()
```

```
[83]: good          50818
      salty         5195
      unknown       1876
      milky          804
      colored        490
      fluoride        217
      Name: quality_group, dtype: int64
```

```
[84]: fig, ax = plt.subplots(figsize=(14,8))
      ax = sns.countplot(x='water_quality', hue="status_group", data=train_data )
      ax.set_title('Water quality with functionality')
```

```
[84]: Text(0.5, 1.0, 'Water quality with functionality')
```



Quantity Column

```
[85]: train_data['quantity'].value_counts()
```


```
[85]: enough          33186
insufficient         15129
dry                  6246
seasonal             4050
unknown              789
Name: quantity, dtype: int64
```

```
[86]: train_data['quantity'].replace('unknown','insufficient',inplace=True)
```

```
[87]: train_data['quantity'].value_counts()
```

```
[87]: enough          33186
insufficient         15918
dry                  6246
seasonal             4050
Name: quantity, dtype: int64
```

```
[88]: fig, ax = plt.subplots(figsize=(10,8))
ax = sns.countplot(x='quantity', hue="status_group", data=train_data)
```



It can be seen obviously that although there are enough water quantity in some wells, they are non-functional. When looking at this graph, dry quantity water points have a highly correlation with non-functionality. If the water point is dry , there is high chance the water point is non functional. On the other hand, if the quantity is enough, there is a higher chance to find functional water points.

```
[89]: train_data['quantity_group'].value_counts()
```

```
[89]: enough          33186
insufficient        15129
dry                 6246
seasonal            4050
unknown              789
Name: quantity_group, dtype: int64
```

```
[90]: train_data['source'].value_counts()
```

```
[90]: spring          17021
shallow well        16824
machine dbh         11075
```

```

river          9612
rainwater harvesting  2295
hand dtw       874
lake           765
dam            656
other          212
unknown        66
Name: source, dtype: int64

```

```

[91]: train_data.loc[train_data['source']
      .value_counts()
      [train_data['source']]
      .values < 900, 'source'] = "Other"

```

```

[92]: train_data['source'].value_counts()

```

```


[92]: spring          17021
      shallow well    16824
      machine dbh     11075
      river           9612
      Other           2573
      rainwater harvesting  2295
      Name: source, dtype: int64

```

```

[93]: fig, ax = plt.subplots(figsize=(10,8))
      ax = sns.countplot(x='source', hue="status_group", data=train_data)

```



Looks like spring waters most valuable according to this graph. If th water point takes spring water mostly water point will be functional. The other water sources almost evenly functional and non functional depend on other features.

```
[94]: train_data['source_type'].value_counts()
```

```
[94]: spring          17021
      shallow well    16824
      borehole        11949
      river/lake      10377
      rainwater harvesting  2295
      dam             656
      other           278
      Name: source_type, dtype: int64
```

```
[95]: train_data['source_class'].value_counts()
```

```
[95]: groundwater    45794
      surface        13328
      unknown         278
```


Name: source_class, dtype: int64

```
[96]: train_data['source_class'].replace('unknown', 'surface', inplace=True)
```

```
[377]: train_data['waterpoint_type'].value_counts()
```


```
[377]: communal standpipe      28522
      hand pump              17488
      other                  6380
      communal standpipe multiple  6103
      improved spring         784
      cattle trough           116
      dam                     7
      Name: waterpoint_type, dtype: int64
```

```
[378]: train_data.loc[train_data['waterpoint_type']
      .value_counts()
      [train_data['waterpoint_type']]
      .values < 900, 'waterpoint_type'] = "other"
```

```
[379]: train_data['waterpoint_type'].value_counts()
```

```
[379]: communal standpipe      28522
      hand pump              17488
      other                  7287
      communal standpipe multiple  6103
      Name: waterpoint_type, dtype: int64
```

```
[385]: fig, ax = plt.subplots(figsize=(12,8))
      ax = sns.countplot(x='waterpoint_type', hue="status_group", data=train_data)
      ax.set_title('Water Point Type with Functionality')
      plt.show()
```



```
[100]: train_data['waterpoint_type_group'].value_counts()
```

```
[100]: communal standpipe    34625
hand pump                    17488
other                       6380
improved spring              784
cattle trough                116
dam                          7
Name: waterpoint_type_group, dtype: int64
```

```
[101]: train_data['longitude'].value_counts()
```

```
[101]: 0.000000    1812
37.540901      2
33.010510      2
39.093484      2
32.972719      2
...
37.579803      1
33.196490      1
34.017119      1
33.788326      1
30.163579      1
```

Name: longitude, Length: 57516, dtype: int64

```
[102]: train_data['longitude'].replace(0.0, train_data
                                         [train_data['longitude'] != 0]
                                         ['longitude'].mean(),
                                         inplace=True)
```

```
[103]: train_data['latitude'].value_counts()
```

```
[103]: -2.000000e-08    1812
       -6.985842e+00     2
       -3.797579e+00     2
       -6.981884e+00     2
       -7.104625e+00     2
       ...
       -5.726001e+00     1
       -9.646831e+00     1
       -8.124530e+00     1
       -2.535985e+00     1
       -2.598965e+00     1
Name: latitude, Length: 57517, dtype: int64
```

```
[104]: train_data[train_data['permit'] == 1]['status_group'].value_counts()
```

```
[104]: functional      21541
       non functional   17311
Name: status_group, dtype: int64
```

```
[105]: train_data[train_data['permit'] == 0]['status_group'].value_counts()
```

```
[105]: functional      10718
       non functional    9830
Name: status_group, dtype: int64
```

```
[106]: # EDA Exploration data analysis
```

```
[107]: le = LabelEncoder()
       train_data['status_group'] = le.fit_transform(train_data['status_group'].values)
```

```
[108]: train_data['status_group'].value_counts()
```

```
[108]: 0    32259
       1    27141
Name: status_group, dtype: int64
```


```
[109]: ## Checking Correlation
       plt.figure(figsize=(12,12))
       corr = train_data.corr()
```

```

mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

sns.heatmap(corr, mask=mask ,annot=True, center=0, vmin=.5, square=True,
            ↳linewidth=.5)
plt.show()

```



0.4.3 Dropping Unnecessary Columns

```

[110]: #Depend on similar values or unnecessary values and multicollinearity , going to
        ↳drop these columns.
train_data.
        ↳drop(columns=['id', 'subvillage', 'payment_type', 'quantity_group', 'num_private',

```

```

        ↳
↳ 'source_type', 'waterpoint_type_group', 'district_code', 'recorded_by',
        ↳
↳ 'scheme_name', 'extraction_type_group', 'lga', 'ward', 'wpt_name'], inplace=True, axis=1)

```

```

[111]: #Checking if we missed something.
train_data.isna().sum()

```

```

[111]: amount_tsh          0
date_recorded             0
funder                    0
gps_height                0
installer                 0
longitude                 0
latitude                  0
basin                     0
region                    0
region_code               0
population                0
public_meeting            0
scheme_management         0
permit                   0
construction_year         0
extraction_type           0
extraction_type_class     0
management                0
management_group          0
payment                   0
water_quality             0
quality_group             0
quantity                  0
source                    0
source_class              0
waterpoint_type           0
status_group              0
dtype: int64

```

```

[112]: # Creating function for metric visualizing.
def get_metrics(clf, X, y):
    y_pred = clf.predict(X)

    my_metrics = (
        (accuracy_score, 'accuracy_score'),
        (recall_score, 'recall_score'),
        (precision_score, 'precision_score'),
        (f1_score, 'f1_score')
    )

```

```

for f, name in my_metrics:
    print(name.title())
    print(f(y, y_pred))
    print()

plot_confusion_matrix(clf, X, y, normalize='true', cmap='Blues')
plt.grid(False)
plt.show()

```

```

[113]: # Importing in-built score modules.
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    classification_report,
    plot_confusion_matrix
)

```

```

[114]: #looking each columns types.
train_data.dtypes

```

```

[114]: amount_tsh          object
date_recorded             object
funder                    object
gps_height                object
installer                 object
longitude                 float64
latitude                  float64
basin                     object
region                    object
region_code               int64
population                object
public_meeting            int64
scheme_management         object
permit                    int64
construction_year         object
extraction_type           object
extraction_type_class     object
management                object
management_group          object
payment                   object
water_quality             object
quality_group             object
quantity                  object
source                    object

```

```
source_class          object
waterpoint_type       object
status_group          int32
dtype: object
```

0.5 First Model(LogReg)

```
[115]: #Create X and y dataframes and train-test split them
y = train_data['status_group']
X = train_data.drop(columns = ['status_group'], axis = 1)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 42)

#Create categorical and continuous feature split
X_train_cat = X_train.select_dtypes('object')
X_train_cont = X_train.select_dtypes(['float64', 'int64'])

#Set up pipeline for scaling continuous variables
continuous_pipeline = Pipeline(steps=[
    ('ss', StandardScaler())
])

#Set up pipeline for encoding categorical variables
categorical_pipeline = Pipeline(steps=[
    ('ohe', OneHotEncoder(handle_unknown='ignore'))
])

#Bind the scaling and encoding process together
transformers = ColumnTransformer(transformers=[
    ('continuous', continuous_pipeline, X_train_cont.columns),
    ('categorical', categorical_pipeline, X_train_cat.columns)
])

#Pipeline for running the model
model1 = Pipeline(steps=[
    ('transformers', transformers),
    ('log', LogisticRegression(class_weight = 'balanced', solver = 'lbfgs',
    ↪random_state=42))
])

#Fitting and checking the score
model1.fit(X_train, y_train)
model1.score(X_train, y_train)
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[115]: 0.7543883277216611
```

```
[116]: model1.steps
```

```
[116]: [('transformers',  
        ColumnTransformer(transformers=[('continuous',  
                                         Pipeline(steps=[('ss', StandardScaler())]),  
                                         Index(['longitude', 'latitude',  
'region_code', 'public_meeting', 'permit'], dtype='object')),  
                                         ('categorical',  
                                          Pipeline(steps=[('ohe',  
OneHotEncoder(handle_unknown='ignore'))])),  
                                         Index(['amount_tsh', 'date_recorded',  
'funder', 'gps_height', 'installer',  
          'basin', 'region', 'population', 'scheme_management',  
          'construction_year', 'extraction_type', 'extraction_type_class',  
          'management', 'management_group', 'payment', 'water_quality',  
          'quality_group', 'quantity', 'source', 'source_class',  
          'waterpoint_type'],  
          dtype='object'))])),  
        ('log', LogisticRegression(class_weight='balanced', random_state=42))]
```

```
[117]: model1[1].coef_.shape
```

```
[117]: (1, 168)
```

```
[118]: # First model score on test data.  
model1.score(X_test, y_test)
```

```
[118]: 0.7521885521885522
```

```
[119]: model1
```

```
[119]: Pipeline(steps=[('transformers',  
                     ColumnTransformer(transformers=[('continuous',  
                                                       Pipeline(steps=[('ss',  
StandardScaler())]),  
                                                       Index(['longitude',  
'latitude', 'region_code', 'public_meeting', 'permit'], dtype='object')),  
                                                       ('categorical',
```



```

Pipeline(steps=[('ohe',
OneHotEncoder(handle_unknown='ignore'))]),
Index(['amount_tsh',
'date_recorded', 'funder', 'gps_height', 'installer',
'basin', 'region', 'population', 'scheme_management',
'construction_year', 'extraction_type', 'extraction_type_class',
'management', 'management_group', 'payment', 'water_quality',
'quality_group', 'quantity', 'source', 'source_class',
'waterpoint_type'],
dtype='object'))]),
('log',
LogisticRegression(class_weight='balanced', random_state=42))])

```

[120]: *# Predicting test data for using on function.*

```

test_pred = model1.predict(X_test)
get_metrics(model1,X_test,y_test)

```

Accuracy_Score

0.7521885521885522

Recall_Score


0.6929798578199052

Precision_Score

0.7443525294304805

F1_Score

0.7177481208774353



```
[121]: # Looking metrics.
print(classification_report(y_test,test_pred))
```

	precision	recall	f1-score	support
0	0.76	0.80	0.78	8098
1	0.74	0.69	0.72	6752
accuracy			0.75	14850
macro avg	0.75	0.75	0.75	14850
weighted avg	0.75	0.75	0.75	14850

```
[122]: # Creating different parameters for gridsearch on forst model(logistic_
      ↪ regression)
ps={
    "log__C":np.logspace(-3,3,7),
    "log__penalty":["l1","l2"]
}
```

```
[123]: # Grid Search on first model.
grids = GridSearchCV(model1,param_grid=ps,cv=5)
```

```
[124]: # Fitting train data.  
       grids.fit(X_train,y_train)
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator  
fit failed. The score on this train-test partition for these parameters will be  
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\pipeline.py", line 335, in fit  
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py", line 1304, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver  
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "  
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator  
fit failed. The score on this train-test partition for these parameters will be  
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\pipeline.py", line 335, in fit  
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py", line 1304, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver  
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "  
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator  
fit failed. The score on this train-test partition for these parameters will be  
set to nan. Details:
```

```
Traceback (most recent call last):
```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit

```

```

        solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit

```

```

    self._final_estimator.fit(Xt, y, **fit_params_last_step)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit

```

```

    self._final_estimator.fit(Xt, y, **fit_params_last_step)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-

```

packages\sklearn\model_selection_validation.py:548: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:

Traceback (most recent call last):

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py", line 531, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\pipeline.py", line 335, in fit
self._final_estimator.fit(Xt, y, **fit_params_last_step)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 1304, in fit
solver = _check_solver(self.solver, self.penalty, self.dual)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 442, in _check_solver
raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py:548: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:

Traceback (most recent call last):

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py", line 531, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\pipeline.py", line 335, in fit
self._final_estimator.fit(Xt, y, **fit_params_last_step)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 1304, in fit
solver = _check_solver(self.solver, self.penalty, self.dual)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 442, in _check_solver
raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py:548: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:

Traceback (most recent call last):

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py", line 531, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-

```

packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "

```

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
```



```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator  
fit failed. The score on this train-test partition for these parameters will be  
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\pipeline.py", line 335, in fit  
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py", line 1304, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver  
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "  
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator  
fit failed. The score on this train-test partition for these parameters will be  
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
```

```

packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)

```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
```

Traceback (most recent call last):

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
```

```
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
```

```
    solver = _check_solver(self.solver, self.penalty, self.dual)
```

```
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
```

```
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
```

```

set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator  
fit failed. The score on this train-test partition for these parameters will be  
set to nan. Details:
```



```

Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

    warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

    warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-

```

```

packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:

```

```

Traceback (most recent call last):

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)

```

```

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-  
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py:548: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\model_selection\_validation.py", line 531, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
```

packages\sklearn\model_selection_validation.py:548: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:

Traceback (most recent call last):

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py", line 531, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\pipeline.py", line 335, in fit
self._final_estimator.fit(Xt, y, **fit_params_last_step)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 1304, in fit
solver = _check_solver(self.solver, self.penalty, self.dual)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 442, in _check_solver
raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py:548: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:

Traceback (most recent call last):

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py", line 531, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\pipeline.py", line 335, in fit
self._final_estimator.fit(Xt, y, **fit_params_last_step)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 1304, in fit
solver = _check_solver(self.solver, self.penalty, self.dual)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py", line 442, in _check_solver
raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py:548: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:

Traceback (most recent call last):

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection_validation.py", line 531, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-

```

packages\sklearn\pipeline.py", line 335, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 1304, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py", line 442, in _check_solver
    raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

```

warnings.warn("Estimator fit failed. The score on this train-test"
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\Users\AI\anaconda3\envs\learn-env\lib\site-
packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[124]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('transformers',
ColumnTransformer(transformers=[('continuous',
Pipeline(steps=[('ss',
                  StandardScaler()))]),
Index(['longitude', 'latitude', 'region_code', 'public_meeting', 'permit'],
dtype='object')),
('categorical',
Pipeline(steps=[('ohe',
                  OneHotEncoder(handle_unknown='ignore'))])),
Index(['amount_tsh', 'date_recorded...
      'construction_year', 'extraction_type', 'extraction_type_class',
      'management', 'management_group', 'payment', 'water_quality',
      'quality_group', 'quantity', 'source', 'source_class',
      'waterpoint_type'],
dtype='object'))]),
      ('log',
```

```
LogisticRegression(class_weight='balanced',
                    random_state=42))),
                    param_grid={'log__C': array([1.e-03, 1.e-02, 1.e-01, 1.e+00,
1.e+01, 1.e+02, 1.e+03]),
                    'log__penalty': ['l1', 'l2']})
```

```
[125]: grids.score(X_train,y_train)
```

```
[125]: 0.7546576879910213
```


```
[126]: # Looking grid search test results for first model.
get_metrics(model1,X_test,y_test)
```

```
Accuracy_Score
0.7521885521885522
```

```
Recall_Score
0.6929798578199052
```

```
Precision_Score
0.7443525294304805
```

```
F1_Score
0.7177481208774353
```



Our first model Logistic Regression didn't give the best value with Grid Search, We are going to forward with other models.

0.6 Second Model(KNN)

```
[127]: # Creating second model pipeline with KNeighbors Classifier.
model2= Pipeline(steps=[
    ('transformers', transformers),
    ('knn', KNeighborsClassifier(n_neighbors=3))
])

#Fitting and checking the score
model2.fit(X_train, y_train)
model2.score(X_train, y_train)
```

```
[127]: 0.8902805836139169
```


```
[128]: # Looking test results for second model.
get_metrics(model2,X_test,y_test)
```

```
Accuracy_Score
0.7938720538720538
```

```
Recall_Score
0.7473341232227488
```

```
Precision_Score
0.7883143258865802
```

```
F1_Score
0.7672774272029196
```



```
[129]: # Creating different parameters for gridsearch on second model(KNeighbors_
      ↪Classifier)
prams = {
    'knn__n_neighbors': [3,5,11,19],
    'knn__weights': ['uniform','distance'],
    'knn__metric':['euclidean','manhattan']
}
```

```
[130]: # Grid Search initializing for second model.
gr = GridSearchCV(model2,param_grid=prams,cv=3)
```

```
[131]: gr.fit(X_train,y_train)
```

```
[131]: GridSearchCV(cv=3,
                    estimator=Pipeline(steps=[('transformers',
ColumnTransformer(transformers=[('continuous',
Pipeline(steps=[('ss',
                    StandardScaler()))])),
Index(['longitude', 'latitude', 'region_code', 'public_meeting', 'permit'],
dtype='object')),
('categorical',
Pipeline(steps=[('ohe',
                    OneHotEncoder(handle_unknown='ignore'))])),
Index(['amount_tsh', 'date_recorded...
```

```

'construction_year', 'extraction_type', 'extraction_type_class',
'management', 'management_group', 'payment', 'water_quality',
'quality_group', 'quantity', 'source', 'source_class',
'waterpoint_type'],
dtype='object'))]])),
                                ('knn',
                                 KNeighborsClassifier(n_neighbors=3))]),
    param_grid={'knn__metric': ['euclidean', 'manhattan'],
                'knn__n_neighbors': [3, 5, 11, 19],
                'knn__weights': ['uniform', 'distance']})

```

```

[132]: # Grid Search train result.
       gr.score(X_train,y_train)

```

```

[132]: 0.997665544332211

```

```

[133]: # Grid Search test results.
       get_metrics(gr,X_test,y_test)

```

```

Accuracy_Score
0.806060606060606

```

```

Recall_Score
0.754739336492891

```

```


Precision_Score
0.8063291139240506

```

```

F1_Score
0.7796817625458997

```



0.7 Third Model(DecisionTree)

```
[134]: # Creating third model pipeline with Decision Tree.
model3= Pipeline(steps=[
    ('transformers', transformers),
    ('tree', DecisionTreeClassifier(criterion='gini',max_depth=5))
])

#Fitting and checking the score
model3.fit(X_train, y_train)
model3.score(X_train, y_train)
```

[134]: 0.7216835016835017

```
[135]: model3
```

```
[135]: Pipeline(steps=[('transformers',
                        ColumnTransformer(transformers=[('continuous',
                                                         Pipeline(steps=[('ss',
                                                                              StandardScaler()))]),
                                                         Index(['longitude',
                                                                              'latitude', 'region_code', 'public_meeting', 'permit'], dtype='object')),
                        ('categorical',
                         Pipeline(steps=[('ohe',
```

```
OneHotEncoder(handle_unknown='ignore'))]),
Index(['amount_tsh',
'date_recorded', 'funder', 'gps_height', 'installer',
'basin', 'region', 'population', 'scheme_management',
'construction_year', 'extraction_type', 'extraction_type_class',
'management', 'management_group', 'payment', 'water_quality',
'quality_group', 'quantity', 'source', 'source_class',
'waterpoint_type'],
dtype='object'))]),
('tree', DecisionTreeClassifier(max_depth=5))])
```


```
[136]: # Looking for metric result for test data.
get_metrics(model3,X_test,y_test)
```

```
Accuracy_Score
0.7188552188552189
```

```
Recall_Score
0.5266587677725119
```

```
Precision_Score
0.7841234840132304
```

```
F1_Score
0.6301054310268451
```



```
[137]: model3.named_steps
```

```
[137]: {'transformers': ColumnTransformer(transformers=[('continuous',  
                                                    Pipeline(steps=[('ss', StandardScaler())]),  
                                                    Index(['longitude', 'latitude', 'region_code',  
'public_meeting', 'permit'], dtype='object')),  
                                                    ('categorical',  
                                                    Pipeline(steps=[('ohe',  
OneHotEncoder(handle_unknown='ignore'))])),  
                                                    Index(['amount_tsh', 'date_recorded',  
'funder', 'gps_height', 'installer',  
'basin', 'region', 'population', 'scheme_management',  
'construction_year', 'extraction_type', 'extraction_type_class',  
'management', 'management_group', 'payment', 'water_quality',  
'quality_group', 'quantity', 'source', 'source_class',  
'waterpoint_type'],  
                                                    dtype='object'))]),  
'tree': DecisionTreeClassifier(max_depth=5)}
```

```
[138]: # Creating different parameters for gridsearch on third model(Decision Tree)  
param = {  
    'tree__criterion': ['gini', 'entropy'],  
    'tree__max_depth': [1,3,5, None],  
    'tree__max_features': ['sqrt', 'log2', None],  
}  
# njobs = -2
```

```
[139]: # Grid Search initializing for third model.  
gr = GridSearchCV(model3,param_grid=param,scoring='recall')
```

```
[140]: # Fitting and looking train set result.  
gr.fit(X_train, y_train)  
gr.score(X_train, y_train)
```

```
[140]: 0.9962234538231399
```


```
[141]: # Looking test set results.  
get_metrics(gr,X_test,y_test)
```

```
Accuracy_Score  
0.7833670033670034
```

```
Recall_Score  
0.7664395734597157
```

```
Precision_Score  
0.7593543653705063
```

F1_Score
0.7628805189061694



```
[142]: gr.best_params_
```

```
[142]: {'tree__criterion': 'gini',  
       'tree__max_depth': None,  
       'tree__max_features': None}
```

```
[143]: gr.best_estimator_
```

```
[143]: Pipeline(steps=[('transformers',  
                      ColumnTransformer(transformers=[('continuous',  
                                                       Pipeline(steps=[('ss',  
                                                           StandardScaler()))]),  
                                                       Index(['longitude',  
                                                           'latitude', 'region_code', 'public_meeting', 'permit'], dtype='object')),  
                                                       ('categorical',  
                                                        Pipeline(steps=[('ohe',  
                                                           OneHotEncoder(handle_unknown='ignore'))])),  
                                                       Index(['amount_tsh',  
                                                           'date_recorded', 'funder', 'gps_height', 'installer',  
                                                           'basin', 'region', 'population', 'scheme_management',
```

```

        'construction_year', 'extraction_type', 'extraction_type_class',
        'management', 'management_group', 'payment', 'water_quality',
        'quality_group', 'quantity', 'source', 'source_class',
        'waterpoint_type'],
        dtype='object')))),
        ('tree', DecisionTreeClassifier()))

```

0.8 Fourth (Ensemble) Model

```

[144]: from sklearn.svm import SVC
       from sklearn.ensemble import BaggingClassifier

```

```

[145]: # Creating fourth model pipeline with ensemble model.
model4= Pipeline(steps=[
    ('transformers', transformers),
    ('bag', BaggingClassifier(XGBClassifier(),
                             n_estimators=10, random_state=0))
])

#Fitting and checking the score
model4.fit(X_train, y_train)
model4.score(X_train, y_train)

```

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[09:28:24] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:25] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:27] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:28] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:30] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the

default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:31] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:32] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:34] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:35] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[09:28:37] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[145]: 0.8503254769921437


[146]: *#Looking test set result.*
get_metrics(model4,X_test,y_test)

Accuracy_Score
0.8117171717171717

Recall_Score
0.7181575829383886

Precision_Score
0.8444792755137582

F1_Score
0.7762125820393788



0.9 Random Forest

```
[147]: # Creating Random Forest model with pipeline.
model5= Pipeline(steps=[
    ('transformers', transformers),
    ('rf', RandomForestClassifier(random_state=42))
])

#Fitting and checking the score
model5.fit(X_train, y_train)
model5.score(X_train, y_train)
```

[147]: 0.9975533108866442


```
[148]: #Looking test set result.
get_metrics(model5,X_test,y_test)
```

Accuracy_Score
0.813063973063973

Recall_Score
0.7711789099526066

Precision_Score
0.8087915501708605

F1_Score
0.7895375284306292



```
[149]: # Looking cross validation for random forest.  
cross_val_score(model5,X_train,y_train)
```

```
[149]: array([0.81661055, 0.81234568, 0.81133558, 0.81661055, 0.80606061])
```

```
[150]: # Creating different paramaters for random forest grid search.  
param_grid = {  
    'rf__n_estimators': [2,5,10,20,50,75,150],  
    'rf__max_features': ['auto', 'sqrt', 'log2'],  
    'rf__max_depth' : [2,5,10,20,50, None],  
    'rf__criterion' :['gini', 'entropy'],  
    'rf__min_samples_split': [2,5,10,20]  
}
```

```
[151]: # CV_rfc = GridSearchCV(estimator=model5, param_grid=param_grid, cv= 5)  
# CV_rfc.fit(X_train, y_train)
```

```
[152]: # CV_rfc.score(X_train,y_train)
```

```
[153]: #get_metrics(CV_rfc,X_test,y_test)
```

```
[154]: # Predicting test set on random forest model.
test_pred = model5.predict(X_test)
```

```
[155]: # Looking metrics.
print(classification_report(y_test, test_pred))
```

	precision	recall	f1-score	support
0	0.82	0.85	0.83	8098
1	0.81	0.77	0.79	6752
accuracy			0.81	14850
macro avg	0.81	0.81	0.81	14850
weighted avg	0.81	0.81	0.81	14850

0.10 Voting Classifier

```
[156]: from sklearn.ensemble import RandomForestClassifier, VotingClassifier
```

```
[157]: # Initializing another ensemble model.
ecclf = VotingClassifier(estimators=[('1', model1), ('2', model2), ('3', model3)])
```

```
[158]: #Fitting train data.
ecclf.fit(X_train, y_train)
```

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[158]: VotingClassifier(estimators=[('1',
                                     Pipeline(steps=[('transformers',
ColumnTransformer(transformers=[('continuous',
Pipeline(steps=[('ss',
                    StandardScaler()))])),
Index(['longitude', 'latitude', 'region_code', 'public_meeting', 'permit'],
dtype='object')),
('categorical',
Pipeline(steps=[('ohe',
OneHotEncoder(handle_unknown='ignore'))])),
```

```

Index(['amount_tsh', 'date_re...
Index(['amount_tsh', 'date_recorded', 'funder', 'gps_height', 'installer',
      'basin', 'region', 'population', 'scheme_management',
      'construction_year', 'extraction_type', 'extraction_type_class',
      'management', 'management_group', 'payment', 'water_quality',
      'quality_group', 'quantity', 'source', 'source_class',
      'waterpoint_type'],
      dtype='object'))]),
      ('tree',
      DecisionTreeClassifier(max_depth=5)))]))

```

```

[159]: #Checking train set score.
      eclf.score(X_train,y_train)

```

```

[159]: 0.8051178451178451

```

```

[160]: #Checking test set score.
      eclf.score(X_test,y_test)

```

```

[160]: 0.7764309764309765

```

```

[161]: # Predicting test set on Voting Classifier.
      preds= eclf.predict(X_test)

```

```

[162]: # Looking metrics.
      print(classification_report(y_test,preds))

```

	precision	recall	f1-score	support
0	0.75	0.88	0.81	8098
1	0.82	0.66	0.73	6752
accuracy			0.78	14850
macro avg	0.78	0.77	0.77	14850
weighted avg	0.78	0.78	0.77	14850

0.11 XGBBoosting with Grid Search

```

[163]: from xgboost import XGBClassifier

```

```

[164]: # Initializing XGB Boosting model.
      model6 = Pipeline(steps=[
          ('transformers', transformers),
          ('xg', XGBClassifier(random_state=42))
      ])

      #Fitting and checking the score

```

```
model6.fit(X_train, y_train)
model6.score(X_train, y_train)
```

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[09:34:43] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[164]: 0.8499438832772166

```
[165]: # Checking test set result.
model6.score(X_test,y_test)
```

[165]: 0.8057912457912458

```
[166]: # Different paramaters for xgb boosting.
param_grid = {
    'xg__learning_rate': [0.1, 0.2],
    'xg__max_depth': [6],
    'xg__min_child_weight': [1, 2],
    'xg__subsample': [0.5, 0.7],
    'xg__n_estimators': [100],
}
```

```
[167]: #Initializing gridsearch and fitting train data.
grid = GridSearchCV(model6,param_grid=param_grid,scoring='recall')
grid.fit(X_train,y_train)
```

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[09:34:44] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:

UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[09:34:46] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[09:34:47] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[09:34:48] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[09:34:49] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:34:51] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:34:52] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:34:53] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:34:55] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```



```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:34:56] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:34:57] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:34:59] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:00] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:01] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:02] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:04] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:05] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:06] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:08] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:09] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:10] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:12] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:13] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:14] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:15] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:17] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:18] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:19] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:21] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:22] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:23] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:25] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:26] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:27] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:28] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:29] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:31] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:32] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:34] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:35] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```

```
C:\Users\AI\anaconda3\envs\learn-env\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated and will be
removed in a future release. To remove this warning, do the following: 1) Pass
option use_label_encoder=False when constructing XGBClassifier object; and 2)
Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[09:35:36] WARNING: ..\src\learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.
```



```
[167]: GridSearchCV(estimator=Pipeline(steps=[('transformers',
ColumnTransformer(transformers=[('continuous',
Pipeline(steps=[('ss',
StandardScaler()))]),
Index(['longitude', 'latitude', 'region_code', 'public_meeting', 'permit'],
dtype='object')),
('categorical',
Pipeline(steps=[('ohe',
OneHotEncoder(handle_unknown='ignore'))])),
Index(['amount_tsh', 'date_recorded', 'fu...

n_estimators=100,
n_jobs=16,
num_parallel_tree=1,
predictor='auto',
random_state=42,
reg_alpha=0, reg_lambda=1,
scale_pos_weight=1,
subsample=1,
tree_method='exact',
validate_parameters=1,
verbosity=None))]),
param_grid={'xg__learning_rate': [0.1, 0.2], 'xg__max_depth': [6],
'xg__min_child_weight': [1, 2],
'xg__n_estimators': [100],
'xg__subsample': [0.5, 0.7]},
scoring='recall')
```

```
[168]: # Grid Search best parameters.
grid.best_params_
```

```
[168]: {'xg__learning_rate': 0.2,
'xg__max_depth': 6,
'xg__min_child_weight': 1,
'xg__n_estimators': 100,
'xg__subsample': 0.5}
```

```
[169]: # Checking train set result.
grid.score(X_train,y_train)
```

```
[169]: 0.7491294325371524
```

```
[170]: # Checking test set result.
grid.score(X_test,y_test)
```

```
[170]: 0.7160841232227488
```


```
[171]: # Looking metrics.
get_metrics(model6,X_test,y_test)
```

Accuracy_Score
0.8057912457912458

Recall_Score
0.7177132701421801

Precision_Score
0.8320741758241759

F1_Score
0.7706743002544529



```
[172]: #Predicting test result.  
preds= grid.predict(X_test)
```

```
[173]: # Checking metrics.  
print(classification_report(y_test,preds))
```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	8098
1	0.83	0.72	0.77	6752
accuracy			0.80	14850

macro avg	0.81	0.80	0.80	14850
weighted avg	0.81	0.80	0.80	14850

0.12 Gradient Boosting

```
[174]: # Initializing Gradient Boosting Classifier with pipeline.
model7 = Pipeline(steps=[
    ('transformers', transformers),
    ('Gbs', GradientBoostingClassifier(random_state=42))
])

#Fitting and checking the score
model7.fit(X_train, y_train)
model7.score(X_train, y_train)
```

```
[174]: 0.770976430976431
```

```
[175]: # Checking test set result.
model7.score(X_test, y_test)
```

```
[175]: 0.7616161616161616
```

```
[176]: # Creating different parameters.
params = {
    'Gbs__learning_rate': [0.075, 0.7],
    'Gbs__max_depth': [13, 14],
    'Gbs__min_samples_leaf': [15, 16],
    'Gbs__max_features': [1.0],
    'Gbs__n_estimators': [100, 200]
}
```

```
[177]: # Initializing grid search and fitting train data.
grid_grad = GridSearchCV(model7, params, cv=5)
grid_grad.fit(X_train, y_train)
```

```
[177]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('transformers',
ColumnTransformer(transformers=[('continuous',
Pipeline(steps=[('ss',
StandardScaler()))]),
Index(['longitude', 'latitude', 'region_code', 'public_meeting', 'permit'],
dtype='object')),
('categorical',
Pipeline(steps=[('ohe',
OneHotEncoder(handle_unknown='ignore'))])),
Index(['amount_tsh', 'date_recorded...
'management', 'management_group', 'payment', 'water_quality',
```

```

        'quality_group', 'quantity', 'source', 'source_class',
        'waterpoint_type'],
        dtype='object'))]])),
        ('Gbs',
        GradientBoostingClassifier(random_state=42))]],
        param_grid={'Gbs__learning_rate': [0.075, 0.7],
                    'Gbs__max_depth': [13, 14], 'Gbs__max_features': [1.0],
                    'Gbs__min_samples_leaf': [15, 16],
                    'Gbs__n_estimators': [100, 200]})

```

```
[193]: grid_grad.best_params_
```

```
[193]: {'Gbs__learning_rate': 0.075,
        'Gbs__max_depth': 14,
        'Gbs__max_features': 1.0,
        'Gbs__min_samples_leaf': 16,
        'Gbs__n_estimators': 200}
```

```
[178]: predss = grid_grad.predict(X_test)
```

```
[179]: #Checking metrics.
print(classification_report(y_test,predss))
```

	precision	recall	f1-score	support
0	0.82	0.88	0.85	8098
1	0.84	0.76	0.80	6752
accuracy			0.83	14850
macro avg	0.83	0.82	0.82	14850
weighted avg	0.83	0.83	0.82	14850


```
[180]: get_metrics(grid_grad,X_test,y_test)
```

```
Accuracy_Score
0.8251851851851851
```

```
Recall_Score
0.7619964454976303
```

```
Precision_Score
0.8387675252689925
```

```
F1_Score
0.7985410523048269
```



```
[402]: cross_val_score(grid_grad,X_test,y_test)
```

```
[402]: array([0.79292929, 0.79427609, 0.79326599, 0.7983165 , 0.78114478])
```

F-1 score explain how good the quality of predictions are and how completely we've predicted labels from dataset . My model predicted %79 percent of data correctly according to f-1 score.

0.12.1 Out of Pipeline model for visualizing features

```
[181]: # Separating numerical and categorical columns.
numerical_cols = X_train.select_dtypes('number').columns.tolist()
categorical_cols = X_train.select_dtypes('object').columns.tolist()
```

```
[182]: # Create objects. (Only processing X-data.)
imputer = SimpleImputer(missing_values=np.nan)
scaler = StandardScaler()
ohe = OneHotEncoder(handle_unknown='ignore', sparse=False)

# Process data.
X_train_num_processed = imputer.fit_transform(X_train[numerical_cols])
X_test_num_processed = imputer.transform(X_test[numerical_cols])

X_train_num_processed = scaler.fit_transform(X_train_num_processed)
X_test_num_processed = scaler.transform(X_test_num_processed)
```

```

X_train_cat_processed = ohe.fit_transform(X_train[categorical_cols])
X_test_cat_processed = ohe.transform(X_test[categorical_cols])

# Join data back together to look at.
X_train_processed_df = pd.DataFrame(
    np.concatenate([X_train_num_processed, X_train_cat_processed], axis=1),
    columns=numerical_cols + ohe.get_feature_names().tolist())
X_test_processed_df = pd.DataFrame(
    np.concatenate([X_test_num_processed, X_test_cat_processed], axis=1),
    columns=numerical_cols + ohe.get_feature_names().tolist())

# Sanity check.
X_train_processed_df.shape, X_test_processed_df.shape

```

[182]: ((44550, 168), (14850, 168))

```

[183]: # Looking first five rows of processed data.
X_train_processed_df.head()

```

```

[183]:  longitude  latitude  region_code  public_meeting  permit  \
0    1.154510 -0.129019   -0.529291      0.408462  0.729442
1   -0.002107  1.940133    0.092332     -2.448208 -1.370910
2   -0.716629  0.983559    0.205354      0.408462  0.729442
3   -0.144681  0.294199   -0.133713      0.408462  0.729442
4   -0.192166 -1.699526   -0.303246     -2.448208  0.729442

    x0_Enough amount water source at this point.  \
0                                           1.0
1                                           0.0
2                                           0.0
3                                           0.0
4                                           0.0

    x0_Less amount water source at this point.  x1_2011  x1_2012  x1_2013  ...  \
0                                           0.0      1.0      0.0      0.0  ...
1                                           1.0      0.0      0.0      1.0  ...
2                                           1.0      1.0      0.0      0.0  ...
3                                           1.0      0.0      0.0      1.0  ...
4                                           1.0      0.0      0.0      1.0  ...

    x18_rainwater harvesting  x18_river  x18_shallow well  x18_spring  \
0                        0.0        0.0                0.0        0.0
1                        0.0        0.0                1.0        0.0
2                        0.0        0.0                0.0        0.0
3                        1.0        0.0                0.0        0.0
4                        0.0        0.0                0.0        1.0

```

	x19_groundwater	x19_surface	x20_communal	standpipe	\
0	1.0	0.0		1.0	
1	1.0	0.0		0.0	
2	1.0	0.0		0.0	
3	0.0	1.0		1.0	
4	1.0	0.0		1.0	

	x20_communal	standpipe	multiple	x20_hand pump	x20_other
0			0.0	0.0	0.0
1			0.0	1.0	0.0
2			0.0	1.0	0.0
3			0.0	0.0	0.0
4			0.0	0.0	0.0

[5 rows x 168 columns]

```
[184]: # Initializing Decision Tree model with best parameters we found at before ↵
↵model.
t = DecisionTreeClassifier(criterion='entropy',max_features=None,max_depth=None)
```


```
[185]: # Fitting train data.
t.fit(X_train_processed_df,y_train)
```

```
[185]: DecisionTreeClassifier(criterion='entropy')
```

```
[186]: # Sorting most important features.
t1 = sorted(t.feature_importances_,reverse=True)
```

```
[398]: # Plot feature importances.

fig, ax = plt.subplots(figsize=(16,8))
for i,r in zip(t1,X_train_processed_df.columns):
    if i > 0.01:
        ax.bar(r, i, color='darkcyan')
        ax.set(title='Feature Importances')
fig.tight_layout()
```



```
[188]: # Initializing Random Forest.
r1 = RandomForestClassifier(random_state=42)
```

```
[189]: # Fitting train data.
r1.fit(X_train_processed_df,y_train)
```

```
[189]: RandomForestClassifier(random_state=42)
```

```
[190]: # Checking random forest most important features.
r1.feature_importances_
```

```
[190]: array([1.61540867e-01, 1.60389667e-01, 1.75112330e-02, 8.84434952e-03,
          9.98999268e-03, 1.01304196e-02, 1.04186628e-02, 4.34926765e-03,
          3.23178127e-03, 4.96438123e-03, 6.60745010e-04, 6.45069865e-04,
          4.87285542e-04, 9.76591752e-05, 2.52943635e-03, 2.70560721e-04,
          6.84564597e-04, 1.27543409e-03, 1.12402140e-03, 2.01194671e-03,
          3.53951168e-04, 1.81330567e-04, 2.03103713e-04, 8.87539011e-03,
          2.03475319e-03, 4.57001049e-04, 5.60687036e-04, 4.21751961e-04,
          4.46247904e-04, 8.72472501e-04, 4.43450165e-04, 9.36279344e-04,
          4.36230330e-04, 4.52979787e-04, 8.92361681e-04, 1.06334789e-02,
          5.60516192e-04, 4.77307033e-04, 1.03690415e-03, 1.95262000e-04,
          3.08581959e-04, 1.71045872e-04, 2.79952280e-04, 1.87109992e-04,
          1.58643353e-03, 1.50712306e-04, 2.00894060e-03, 7.96877358e-04,
          1.35530898e-03, 3.34810062e-04, 2.55275220e-03, 6.83559427e-04,
          5.61036532e-04, 2.21554015e-03, 1.52109469e-03, 7.58630083e-04,
          3.78364642e-03, 4.24085501e-03, 9.07284971e-03, 8.46286212e-03,
          8.59049539e-03, 5.04213401e-03, 3.34538282e-03, 2.74786166e-03,
          3.23719518e-03, 2.98825724e-03, 3.09044719e-03, 3.15193760e-03,
```



```

2.28321289e-03, 2.90408846e-03, 2.62123900e-03, 3.55431260e-04,
1.30893920e-03, 5.99323089e-03, 1.49486820e-03, 2.49925989e-03,
2.00704939e-03, 1.07287166e-03, 1.23104966e-03, 1.43115422e-03,
1.41267137e-03, 1.87615531e-03, 1.26627821e-03, 1.94317124e-03,
1.51139980e-03, 1.22838807e-03, 9.96265059e-04, 1.72518929e-03,
1.52262086e-03, 9.41154481e-04, 1.57574196e-03, 8.33305881e-03,
7.53913630e-03, 7.89429431e-03, 1.22651811e-03, 3.05699029e-03,
1.15043102e-03, 1.16841733e-03, 1.04029727e-03, 6.12731339e-03,
1.69068073e-03, 2.56484129e-03, 2.32822912e-03, 2.78773440e-03,
9.29957831e-03, 8.72709239e-03, 1.16400956e-03, 7.94339858e-03,
5.90639113e-03, 6.33536222e-03, 2.37648114e-03, 2.47327228e-03,
4.96268258e-03, 2.66441076e-03, 1.84228627e-03, 2.52760180e-03,
6.14361610e-03, 1.77981404e-02, 2.81462215e-03, 3.58518794e-03,
4.10281336e-03, 6.04098434e-03, 2.35707146e-03, 1.38022283e-02,
3.18208296e-03, 2.08192238e-03, 2.14095956e-03, 6.88253164e-03,
1.74872037e-03, 2.24737839e-03, 1.49458735e-03, 2.98330027e-03,
2.67045583e-03, 1.00847739e-03, 1.38545097e-03, 6.90219412e-04,
3.90772584e-03, 1.33228429e-02, 1.22157597e-03, 3.24768303e-03,
4.86284434e-03, 7.03860312e-03, 3.47527997e-03, 5.52534350e-03,
4.12326542e-03, 4.02016517e-03, 7.92005770e-04, 4.76629972e-04,
4.09946731e-03, 1.08534386e-03, 3.26326785e-03, 4.94810098e-03,
6.49035737e-02, 3.07653746e-02, 1.34818225e-02, 5.97374714e-03,
3.02912270e-03, 5.04951606e-03, 3.16213180e-03, 3.44944763e-03,
5.34826047e-03, 6.60568584e-03, 4.78113837e-03, 4.48058073e-03,
1.25003780e-02, 8.79952043e-03, 7.10986558e-03, 2.08018851e-02])

```

```


[191]: # Sorting important features.
r11 = sorted(r1.feature_importances_,reverse=True)

```

```

[401]: # Plot feature importances.
fig, ax = plt.subplots(figsize=(12,8))
for i,r in zip(r1.feature_importances_,X_train_processed_df.columns):
    if i > 0.018:
        ax.bar(r, i, color='darkcyan')
        ax.set(title='Feature Importances')
fig.tight_layout()

```



0.13 Conclusion

In conclusion , built model is predictive of functionality of water wells in Tanzania with a **F-1 Score** 0.79. Validated this score with train test split and cross validation. Final model included 168 variables, most of them one hot encoded columns.

The main metric that I would be using to assess my models' performance here is F-1 Score. F-1 score explain how good the quality of predictions are and how completely we've predicted labels from dataset. We wouldn't look at accuracy score because it would be misleading for our specific project. Because accuracy generally good for balanced classes and if both classes importances the same. We are goin to look at F-1 Score because it is harmonic mean of precision and recall scores what exactly need for this project. Which is for this project 0 Non-Functional class important for us.

0.14 Future Work

1. Gather better quality data for prediction model.
2. Bring together old and new data for preparing for modeling.
3. Work on models to predict better.